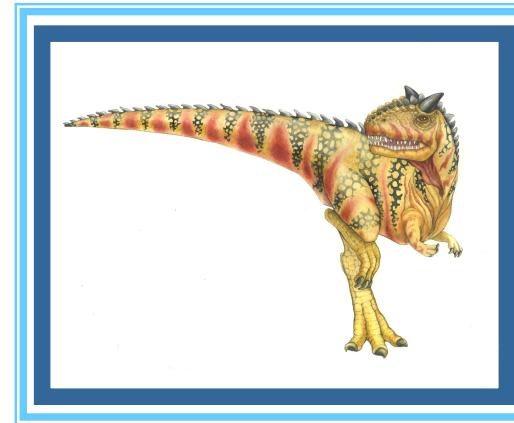


Memory Management



Slides mostly borrowed from Galvin, with slight modifications by us



Memory Management: Topics

- Background
- Contiguous Memory Allocation
- Segmentation
- Paging
- Structure of the Page Table





Multiple processes share the memory

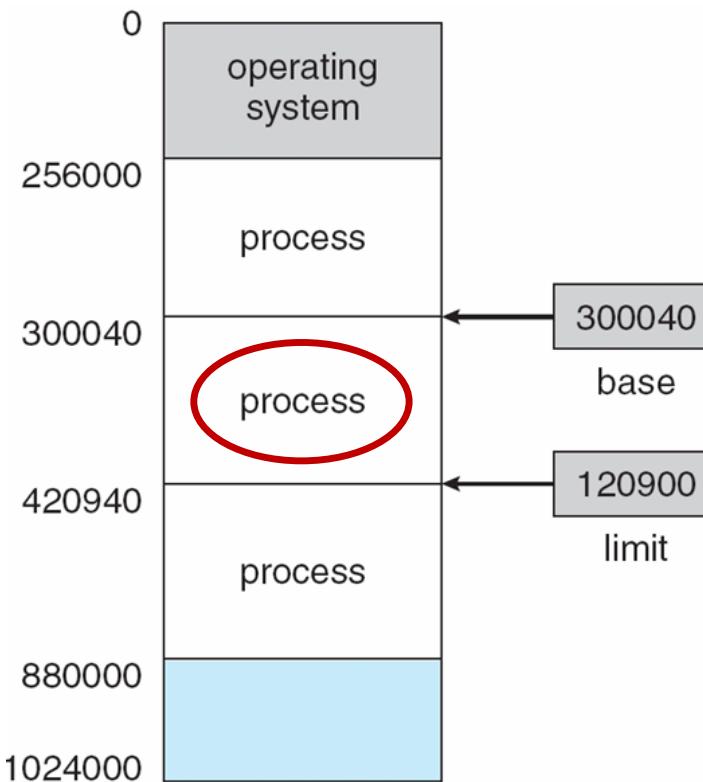
- By "Memory", we refer to "Main memory"
- Program must be brought (from disk) into memory and placed within a process for it to be run
- Multiple programs (processes) share the memory
- Protection of memory required to ensure correct operation of various processes that share the memory (see next slide)





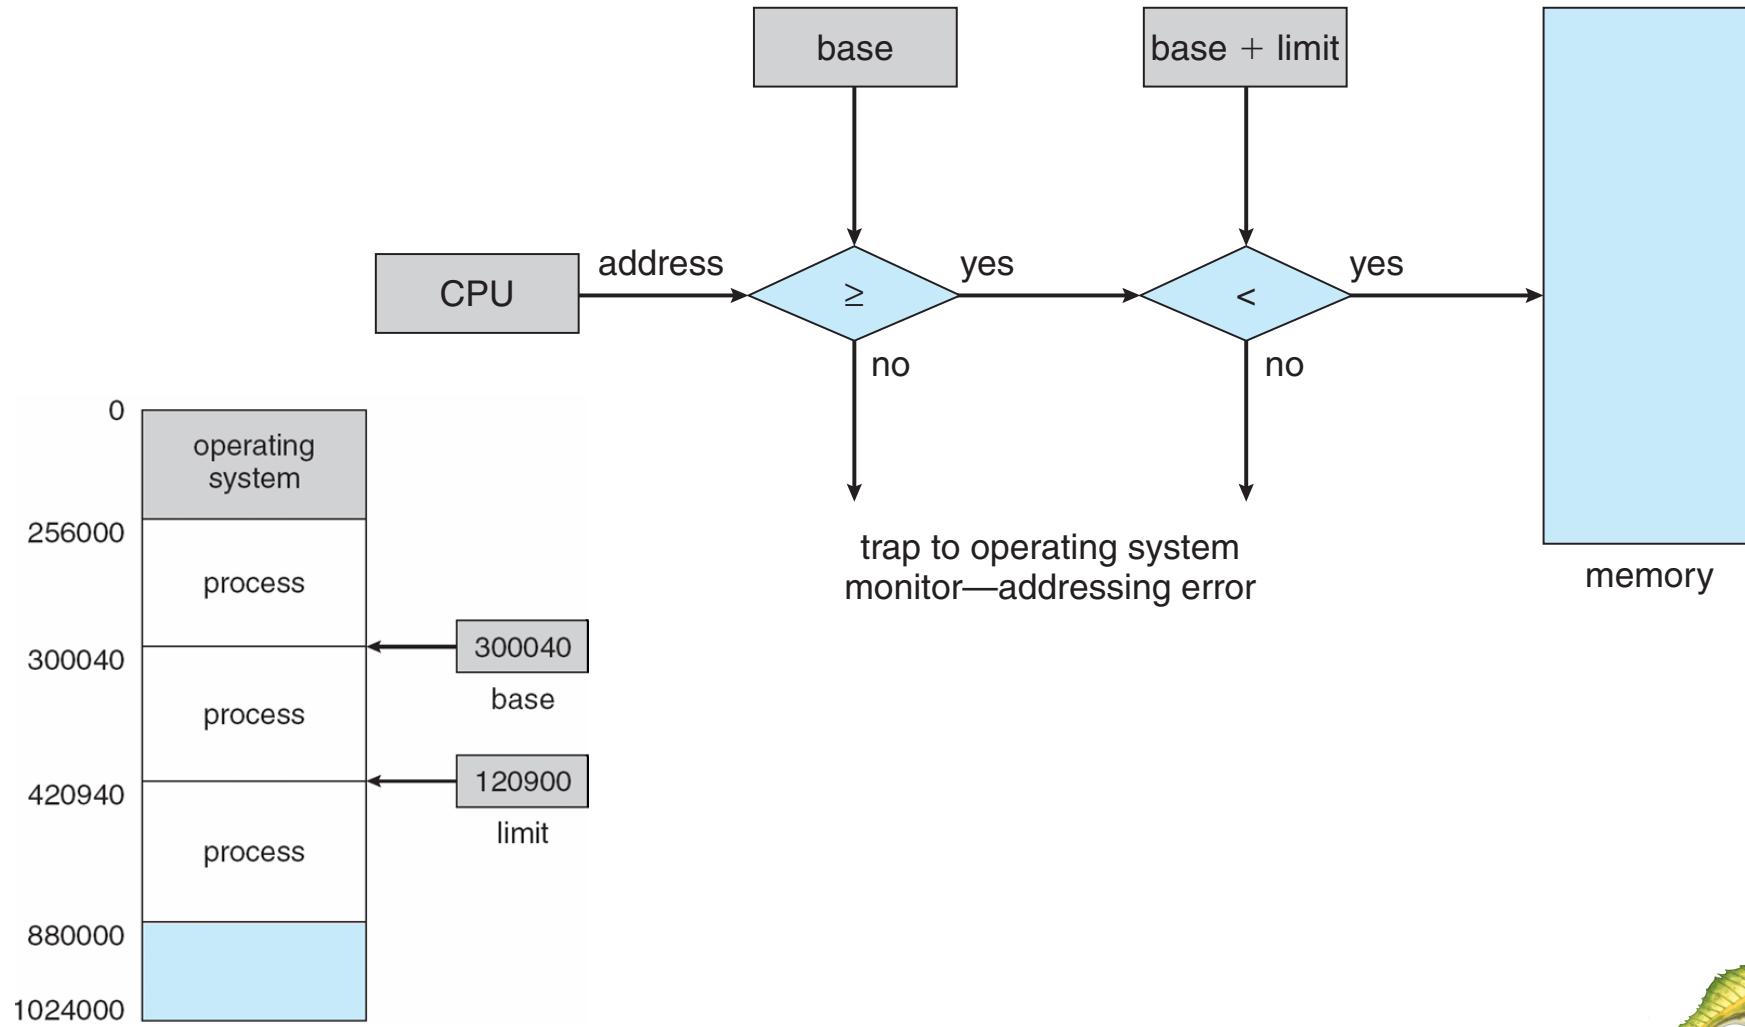
Base and Limit Registers

- A pair of **base** and **limit registers** define the logical address space of a user process
- CPU must check every memory access generated by a process in user mode to be sure it is between base and limit for that process





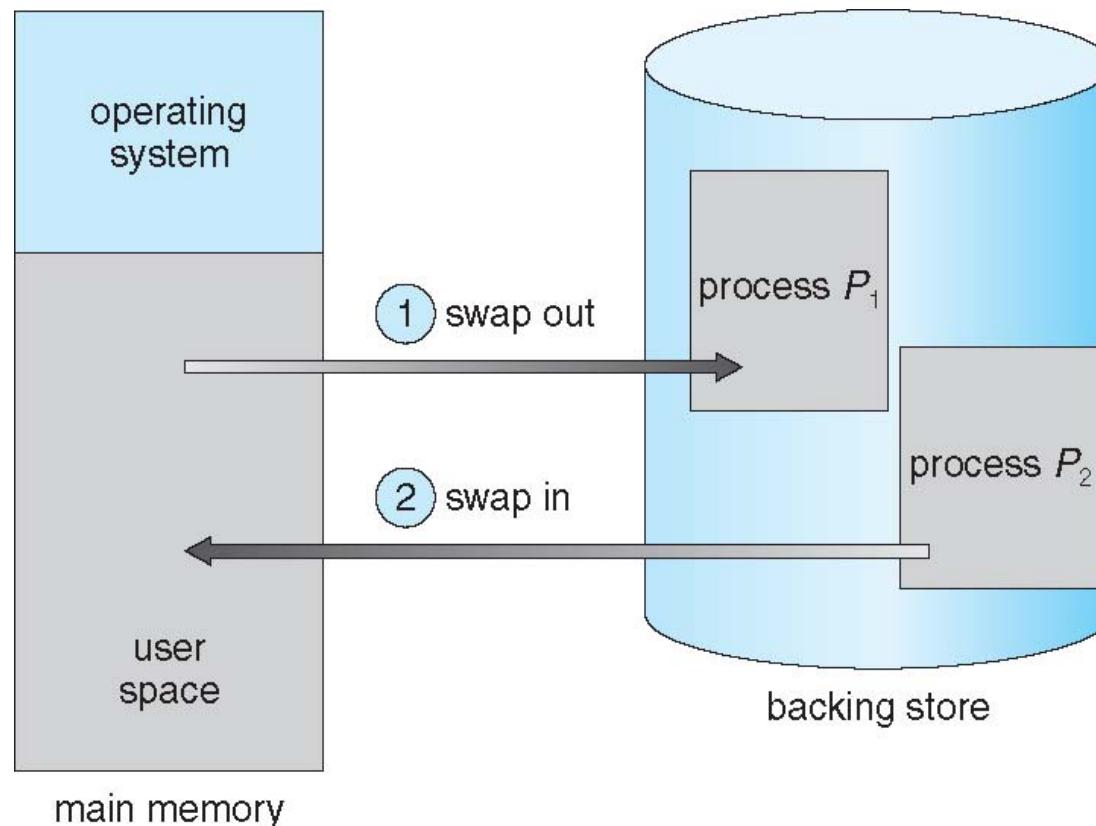
Hardware Address Protection





Swapping

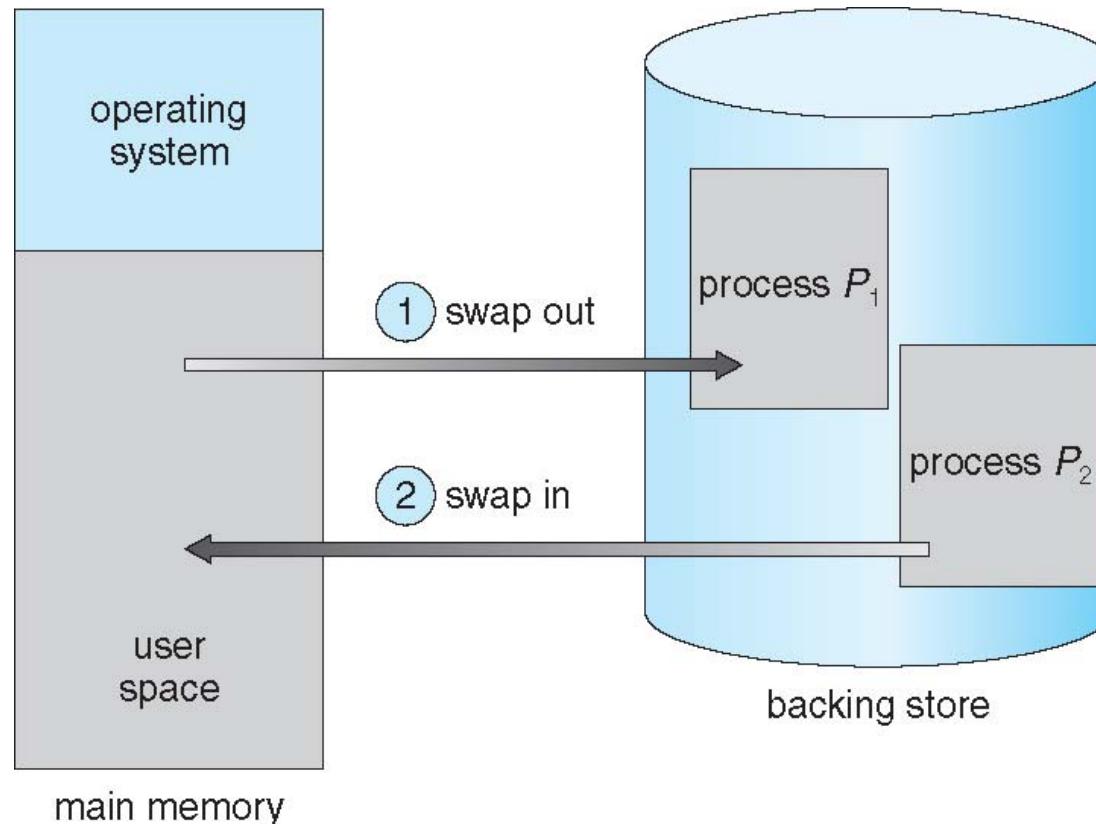
- A process can be **swapped** temporarily out of memory to a backing store, and then brought back into memory for continued execution
- **Backing store** – fast disk large enough to accommodate copies of all memory images for all user processes





Swapping

- A process can be **swapped** temporarily out of memory to a backing store, and then brought back into memory for continued execution
- **Backing store** – fast disk large enough to accommodate copies of all memory images for all user processes



Note: Total physical memory space of all running processes can exceed the size of physical memory





What a programmer wants from memory

- Memory addresses will start from 0
- Memory will be contiguous
- I should be able to dynamically increase the memory
- I should have as much memory as I want

In systems where multiple programs (processes) are in memory at the same time, EACH program has the above expectations





Address Binding

- We want to write programs without worrying about where in memory our program will be loaded during execution (we want to assume that memory will start from address zero)
 - Need a mechanism by which user processes can be stored in different areas of the memory
 - Need mechanisms to map addresses generated by a user program/process to actual physical memory addresses
- Address Binding: mapping addresses from one address space to another





Logical vs. Physical Address Space

- The concept of a **logical address space** that is bound to a separate **physical address space** is central to proper memory management
- **Logical address** – generated by the CPU; also referred to as **virtual address**
 - **Logical address space** is the set of all logical addresses generated by a program (*i.e., the CPU*)
- **Physical address** – address seen by the memory unit
 - **Physical address space** is the set of all physical addresses generated by a program (*i.e., accessed in main memory*)





Memory-Management Unit (MMU)

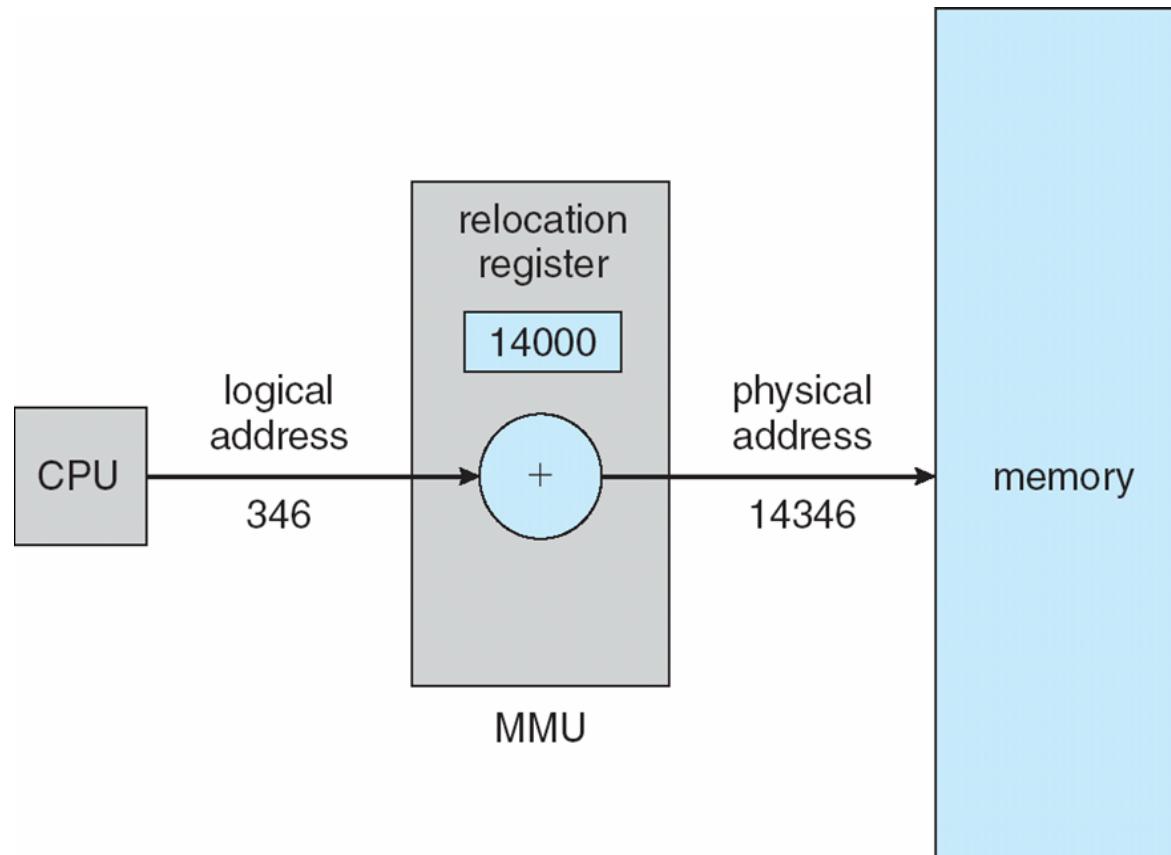
- Hardware device that at run time maps virtual to physical address
- The user program deals with *logical* addresses; it never sees the *real* physical addresses
 - Execution-time binding occurs when reference is made to location in memory
 - Logical address bound to physical addresses





Dynamic relocation using a relocation register

- Consider a simple scheme where the value in a **relocation register** is added to every address generated by a user process at the time it is sent to memory





What a programmer wants from memory

- Memory addresses will start from 0 Satisfied:
Relocation register
- Memory will be contiguous
- I should be able to dynamically increase the memory
- I should have as much memory as I want

In systems where multiple programs (processes) are in memory at the same time, EACH program has the above expectations





Memory Management: Topics

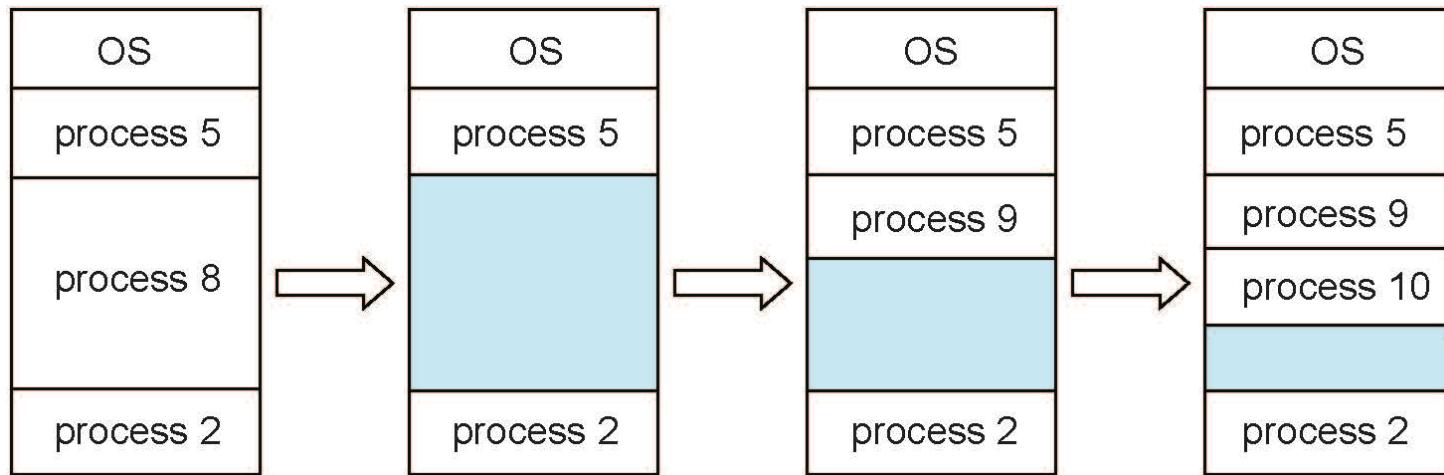
- Background
- Contiguous Memory Allocation
- Segmentation
- Paging
- Structure of the Page Table





Contiguous Allocation

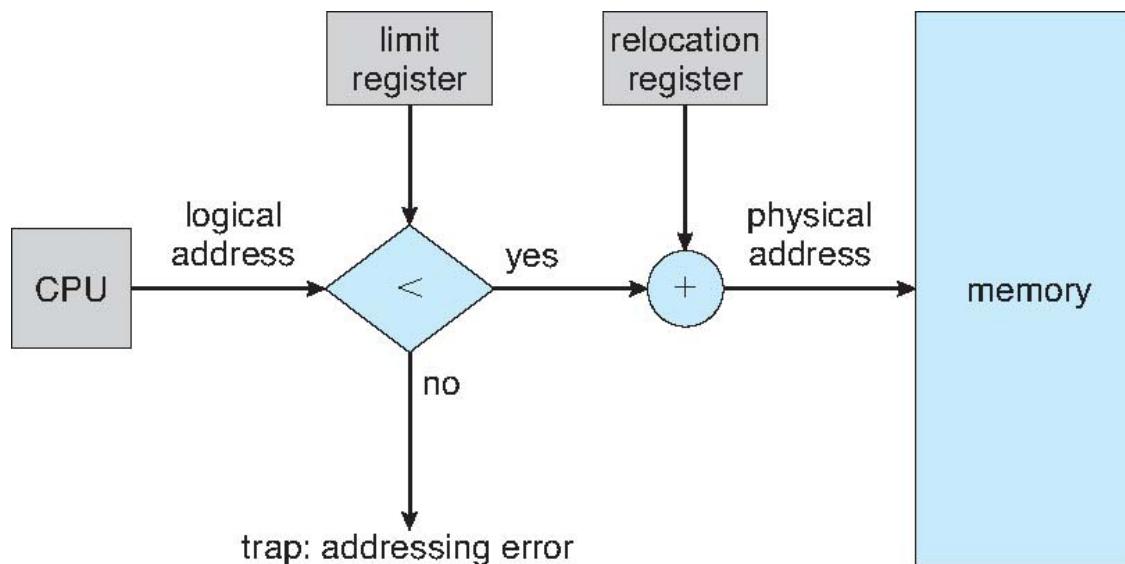
- Main memory must support both OS and user processes
- Limited resource, must allocate efficiently
- **Contiguous allocation** is one early method
- Main memory usually is split into two partitions:
 - a) Resident OS, usually held in lower address portion of memory
 - b) User processes then held in higher address portion of memory
 - ▶ Each process contained in single contiguous section of memory





Contiguous Allocation (Cont.)

- Relocation registers used to protect user processes from each other, and from changing operating-system code and data
 - **Base / relocation register** contains value of smallest physical address of a user process
 - **Limit register** contains range of logical addresses accessible by a user process – each logical address must be less than the limit register
 - MMU maps logical address *dynamically*
 - Instructions to load these registers must be **privileged**

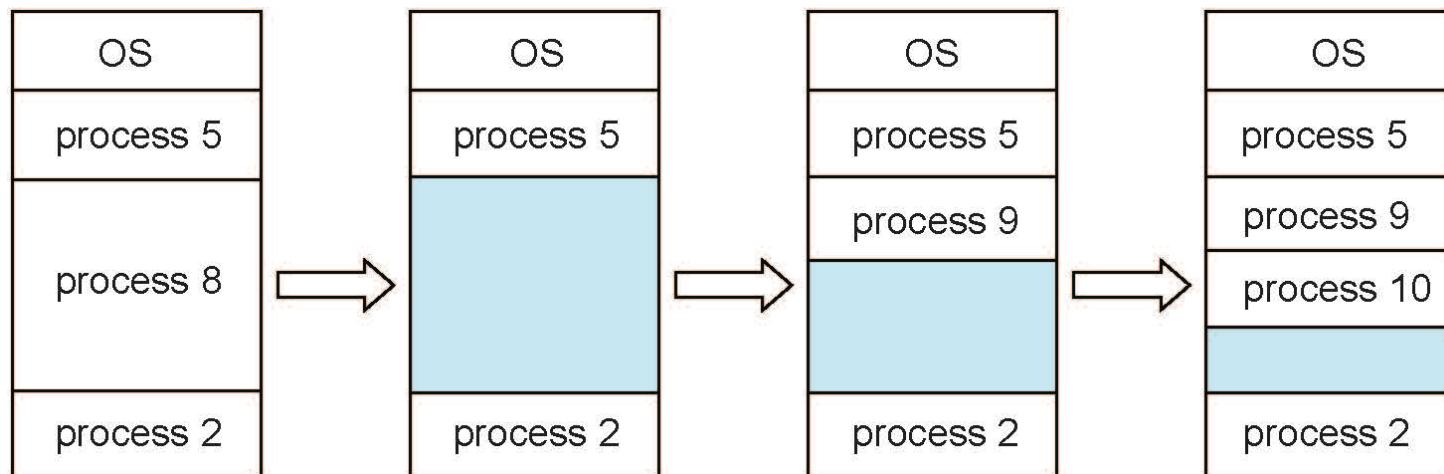




Multiple-partition allocation

■ Multiple-partition allocation

- **Variable-size partitions** (sized to a given process' needs)
- **Hole** – block of available memory; holes of various size are scattered throughout memory
- When a process arrives, it is allocated memory from a hole large enough to accommodate it
- Process exiting frees its partition, adjacent free partitions combined
- Operating system maintains information about:
 - allocated partitions
 - free partitions (hole)





Dynamic Storage-Allocation Problem

How to satisfy a request of size n from a list of free holes?

- **First-fit:** Allocate the *first* hole that is big enough
- **Best-fit:** Allocate the *smallest* hole that is big enough
 - Benefit: Produces the smallest leftover hole
 - Cost: Must search entire list, unless ordered by size
- **Worst-fit:** Allocate the *largest* hole
 - Must also search entire list
 - Produces the largest leftover hole

First-fit and best-fit better than worst-fit in terms of speed and storage utilization





Fragmentation – a problem

- **External Fragmentation** – total memory space exists to satisfy a request, but it is not contiguous
 - Holes of sizes 10K, 15K and 25K exists
 - A process of size 35K arrives --- cannot be loaded

- **Internal Fragmentation** – allocated memory may be slightly larger than requested memory; this size difference is memory internal to a partition, but not being used
 - Allocated memory = 20K, but using only 18K
 - 2K wasted due to internal fragmentation





Fragmentation (Cont.)

- Reduce external fragmentation by **compaction**
 - Shuffle memory contents to place all free memory (i.e. holes) together in one large block
 - Compaction is usually very expensive
 - Compaction is possible *only* if relocation is dynamic, and is done at execution time
- For the same reasons, dynamically increasing memory of a process is very expensive





What a programmer wants from memory

- Memory addresses will start from 0 Satisfied:
Relocation register
- Memory will be contiguous Easy with contiguous allocation, but fragmentation, wastage of memory. Dynamic increase of mem impractical
- I should be able to dynamically increase the memory
- I should have as much memory as I want

In systems where multiple programs (processes) are in memory at the same time, EACH program has the above expectations





Memory Management: Topics

- Background
- Contiguous Memory Allocation
- Segmentation
- Paging
- Structure of the Page Table

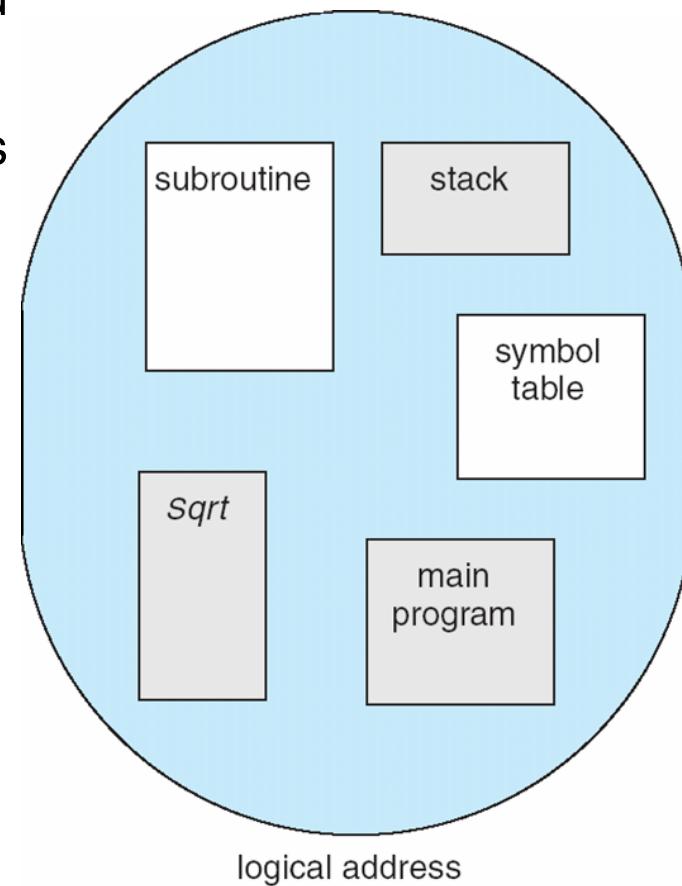




Segmentation

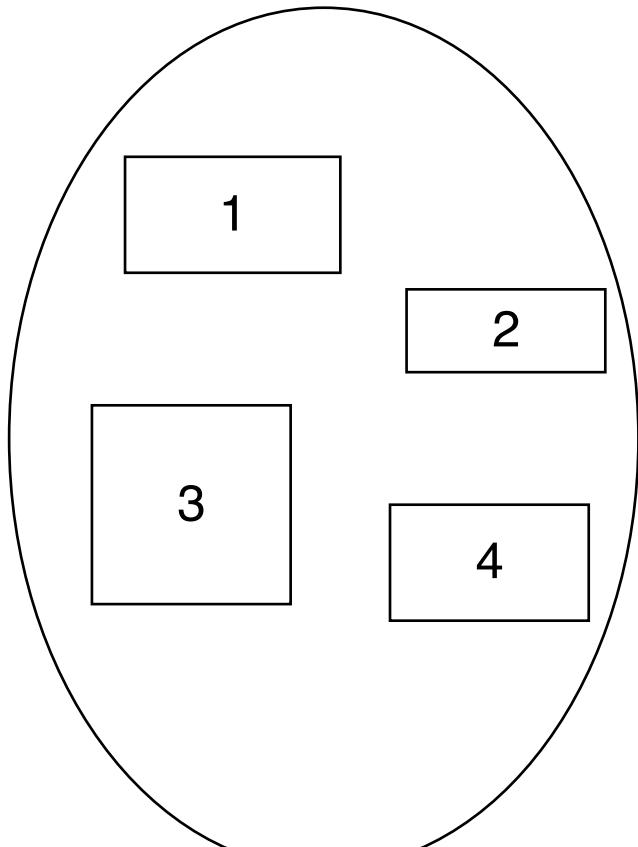
- Memory-management scheme that supports user-view of memory
- A program is a collection of segments, where a segment is a logical unit such as:
 - a function / procedure / method
 - an object
 - local variables, global variables
 - stack

User's view of a program

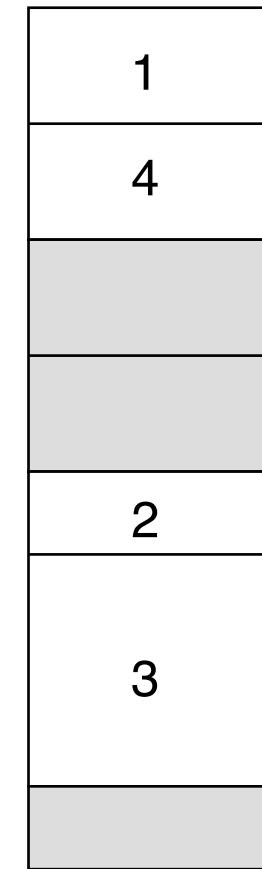




Logical View of Segmentation



user space



physical memory space

Note: we are shifting to non-contiguous memory allocation





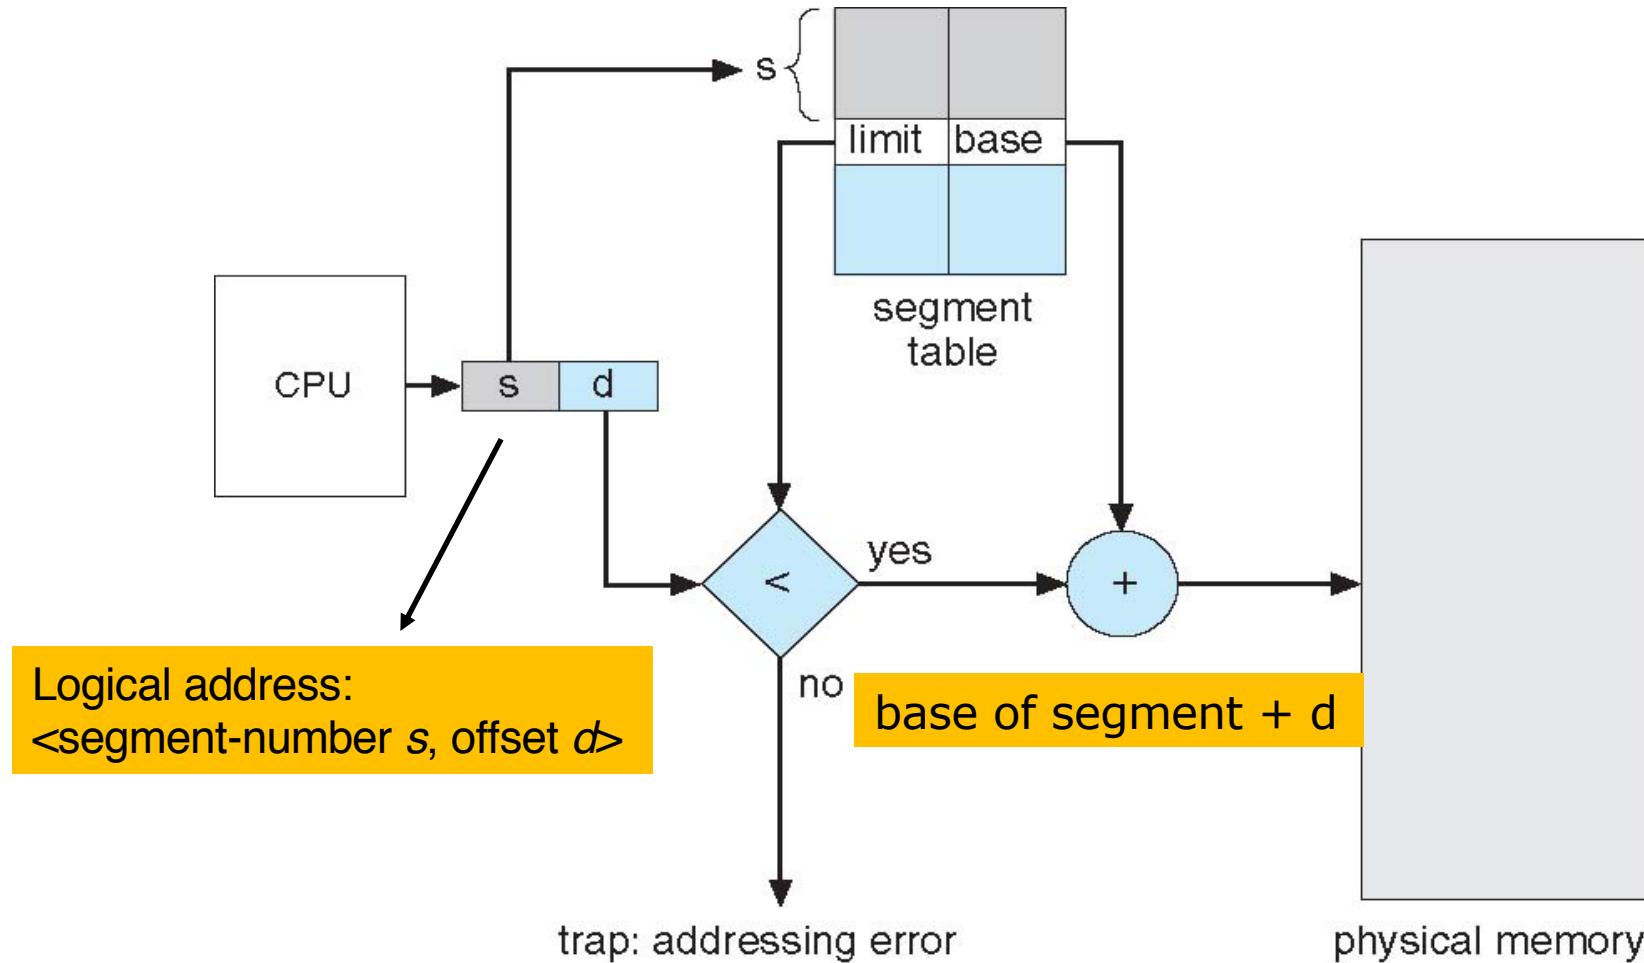
Segmentation Architecture

- Logical address consists of a two-tuple:
 $\langle \text{segment-number } s, \text{ offset } d \rangle$
- **Segment table** – used to map logical addresses to physical addresses
- **Segment table has one entry for each segment of this process**
- Each segment table entry has:
 - **base** – contains the starting physical address where the corresponding segment resides in memory
 - **limit** – specifies the length of the segment (d should be less than $limit$)
 - Other fields, such as protection bits, access privileges





Segmentation Hardware





Memory Management: Topics

- Background
- Contiguous Memory Allocation
- Segmentation
- **Paging**
- Structure of the Page Table





Paging

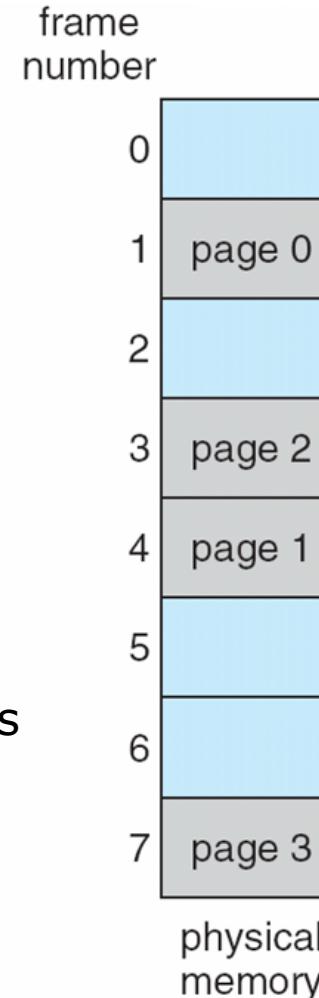
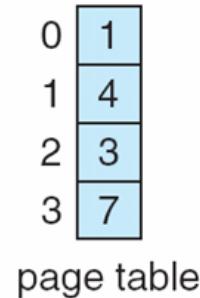
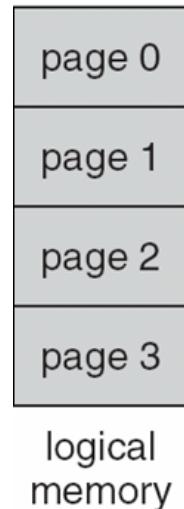
- Divide physical memory into **fixed-sized blocks** called **frames**
 - Size is power of 2, between 512 bytes and 16 MB
 - Typical size: 32KB, 4 MB
- Divide logical memory into blocks of same size called **pages**
- Physical address space of a process can be noncontiguous; **allocated physical memory in units of a page wherever available** (see next slide)





Paging Model of Logical and Physical Memory

Program's view:
contiguous



Physically
non-contiguous
allocation

To run a program of size N pages, OS needs to find N free frames and load program (frames need not be contiguous)

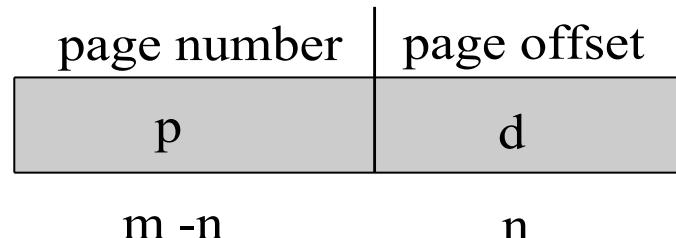
Set up a **page table** to translate logical to physical addresses





Address Translation Scheme

- Address generated by CPU is divided into:
 - **Page number (p)** – used as an index into a **page table** which contains base address of each page in physical memory
 - **Page offset (d)** – combined with base address to define the physical memory address that is sent to the memory unit

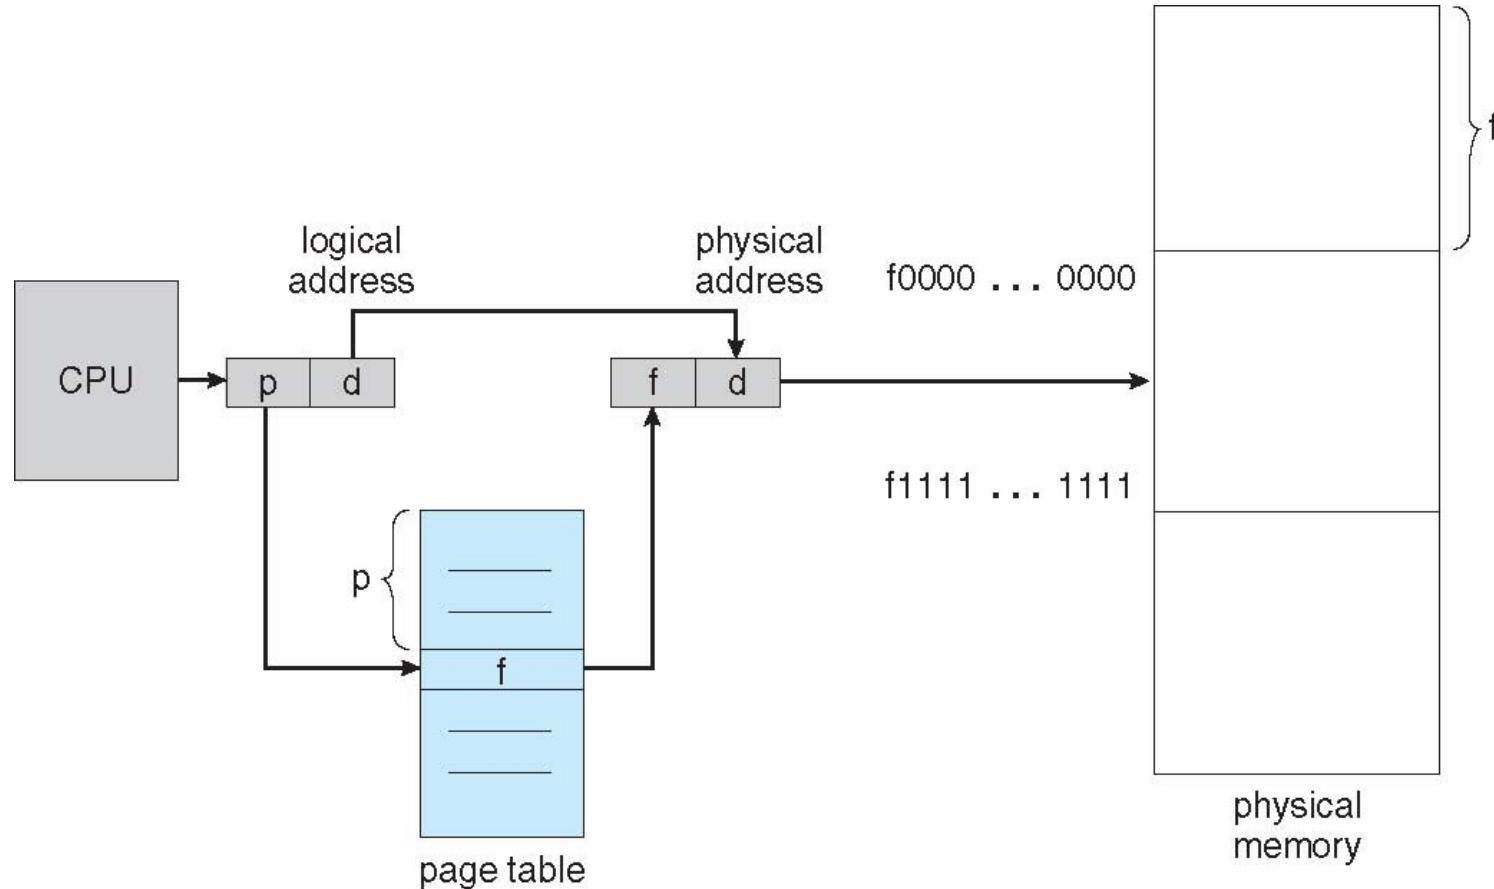


- For given logical address space 2^m and page size 2^n





Paging Hardware





Paging Example

Program's view:
contiguous

0	a
1	b
2	c
3	d
4	e
5	f
6	g
7	h
8	i
9	j
10	k
11	l
12	m
13	n
14	o
15	p

logical memory

0	5
1	6
2	1
3	2

page table

0	
4	i j k l
8	m n o p
12	
16	
20	a b c d
24	e f g h
28	

physical memory

page number	page offset
p	d

m - n n

n=2 and m=4

n=2 means 4-byte pages

32-byte physical memory = 8 frames

$2^4 = 16$ -byte logical address space

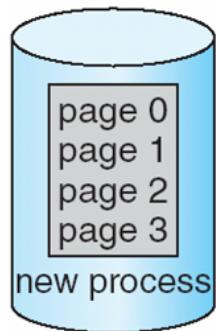




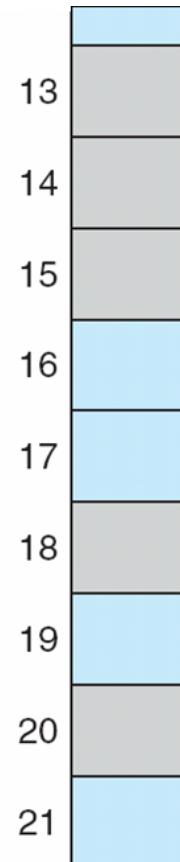
Free Frames

free-frame list

14
13
18
20
15



OS maintains info of
free frames
(frame table)

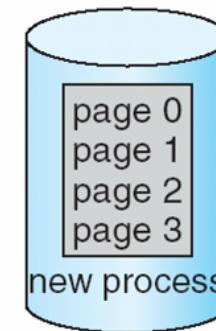


(a)

Before allocation

free-frame list

15



new-process page table

0	14
1	13
2	18
3	20

(b)

After allocation





How to decide page / frame size?

■ Calculating internal fragmentation

- Page size = 2048 bytes
- Process size = 72,766 bytes $(35 * 2048 + 1086)$
- 35 pages + 1086 bytes
- Internal fragmentation of $2048 - 1086 = 962$ bytes

■ How to decide frame or page size?

- Worst case fragmentation = 1 frame – 1 byte
- Average fragmentation = 1/2 of frame size
- So small frame sizes desirable?
- But each page table entry takes memory to track





Optimum Page Size

- Assume that:
 - Average process size = S
 - Page table entry size = K bytes
 - Page size = P bytes
- Average internal fragmentation per process = $P / 2$
- Average number of pages per process = S / P (actually ceiling)
- Thus, total overhead $V = KS / P + P / 2$
- To find the value of P that minimizes overhead, set $dV/dP = 0$
 - $-KS / P^2 + 1/2 = 0$
 - Thus, $P = \sqrt{2SK}$

Small P implies less internal fragmentation

Large P implies smaller page table (lower overhead)





Memory Protection in Paging

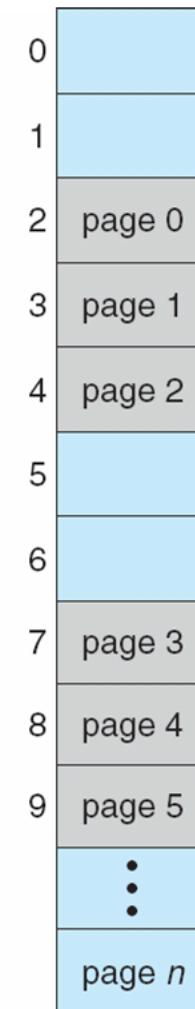
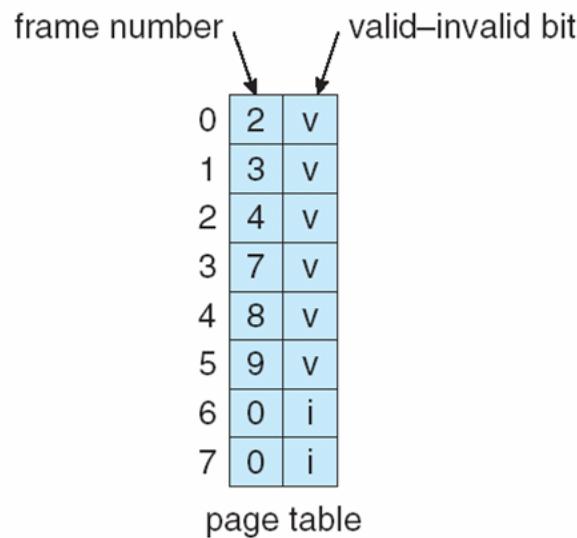
- Memory protection implemented by associating **protection bits** with each page (these bits usually kept in the page table)
 - Indicates if read-only or read-write access is allowed
 - Can also add more bits to indicate page execute-only, and so on
- **Valid-invalid** bit attached to each entry in the page table:
 - “valid” indicates that the associated page is in the process’ logical address space, and is thus a legal (valid) page
 - “invalid” indicates that the page is not in the process’ logical address space
- Any violations result in a trap to the kernel





Valid (v) or Invalid (i) Bit In A Page Table

00000	page 0
	page 1
	page 2
	page 3
	page 4
10,468	page 5
12,287	





Advantages of Paging

- Clear separation between the programmer's view of memory and the actual physical memory
- Programmer's view of memory of his/her process
 - A single contiguous space, starting at (logical) address 0
 - Assume entire physical memory contains only this process
- Actual physical memory
 - Non-contiguous, scattered in pages across the physical memory
 - Occupies only a part of the physical memory
- Avoids external fragmentation, but internal fragmentation can still exist





Advantages of paging

■ Shared code

- One copy of read-only code shared among multiple processes (i.e., multiple processes running a text editor during a lab class)
- Also useful for inter-process communication if sharing of read-write pages is allowed

■ Private code and data

- Each process might also keep its own copy of private code / data
- The pages for the private code / data can appear anywhere in the logical address space

■ Dynamic increase of memory

- If a process needs more memory, can allocate an empty frame as an additional page





What a programmer wants from memory

- Memory addresses will start from 0
- Memory will be contiguous First 3 requirements satisfied with paging
- I should be able to dynamically increase the memory
- I should have as much memory as I want Need to satisfy requirements for large memory

In systems where multiple programs (processes) are in memory at the same time, EACH program has the above expectations





Memory Management: Topics

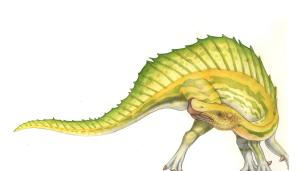
- Background
- Contiguous Memory Allocation
- Segmentation
- Paging
- Structure of the Page Table (and variants of the page table) – dealing with large Page Tables efficiently





Implementation of Page Table

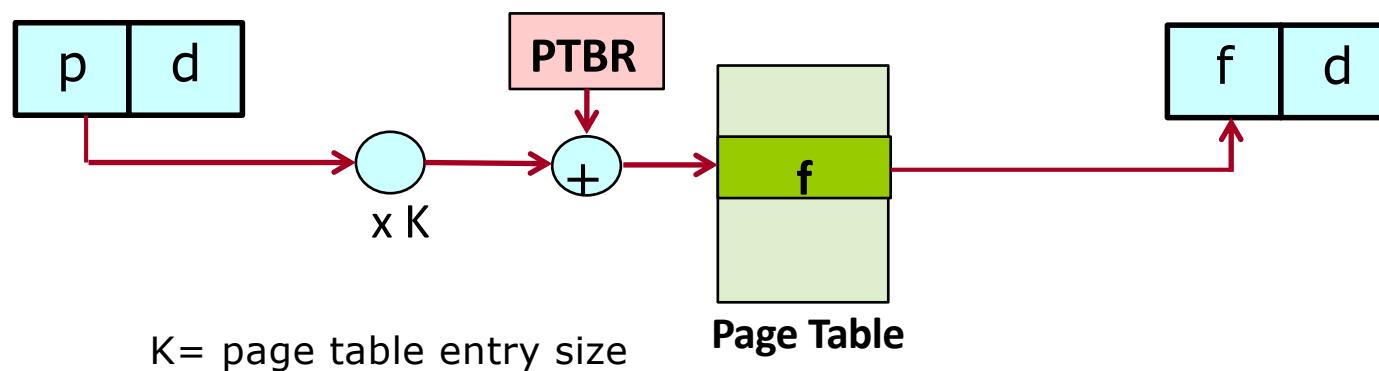
- Page Table has to be accessed for every memory access by a process
 - Implementation should allow fast access
- Where to store the Page Table?
- Best case – a set of dedicated registers
 - Only feasible if the page table is small
 - Not practical in most modern systems
- Practical case – **Page Table stored in the memory**





Implementation of Page Table

- In practice, page table is kept in main memory
- **Page-table base register (PTBR)** points to the page table
 - In this scheme **every data/instruction access requires two memory accesses**
 - One for the page table entry, and one for the actual data/instruction
 - Too expensive





Implementation of Page Table (Cont.)

- The two memory access problem can be solved by the use of a special fast-lookup **hardware cache** called **associative memory** or **translation look-aside buffers (TLBs)**
- TLBs typically small (between 32 and 1,024 entries)
 - Contains a few of the (most frequently/recently accessed) page table entries
 - For a logical address generated by CPU, its page number is searched in TLB; If page number is found (**TLB hit**), its frame number is immediately available to access the data directly in memory
- On a **TLB miss**, a memory reference to the page table needed
 - Value is loaded into the TLB for faster access next time
 - Replacement policies must be considered
 - Some entries can be **wired down** for permanent fast access



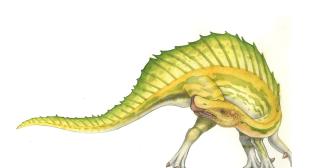


Associative Memory

- TLB is associative memory – allows parallel search (item to be searched is compared with all keys simultaneously)

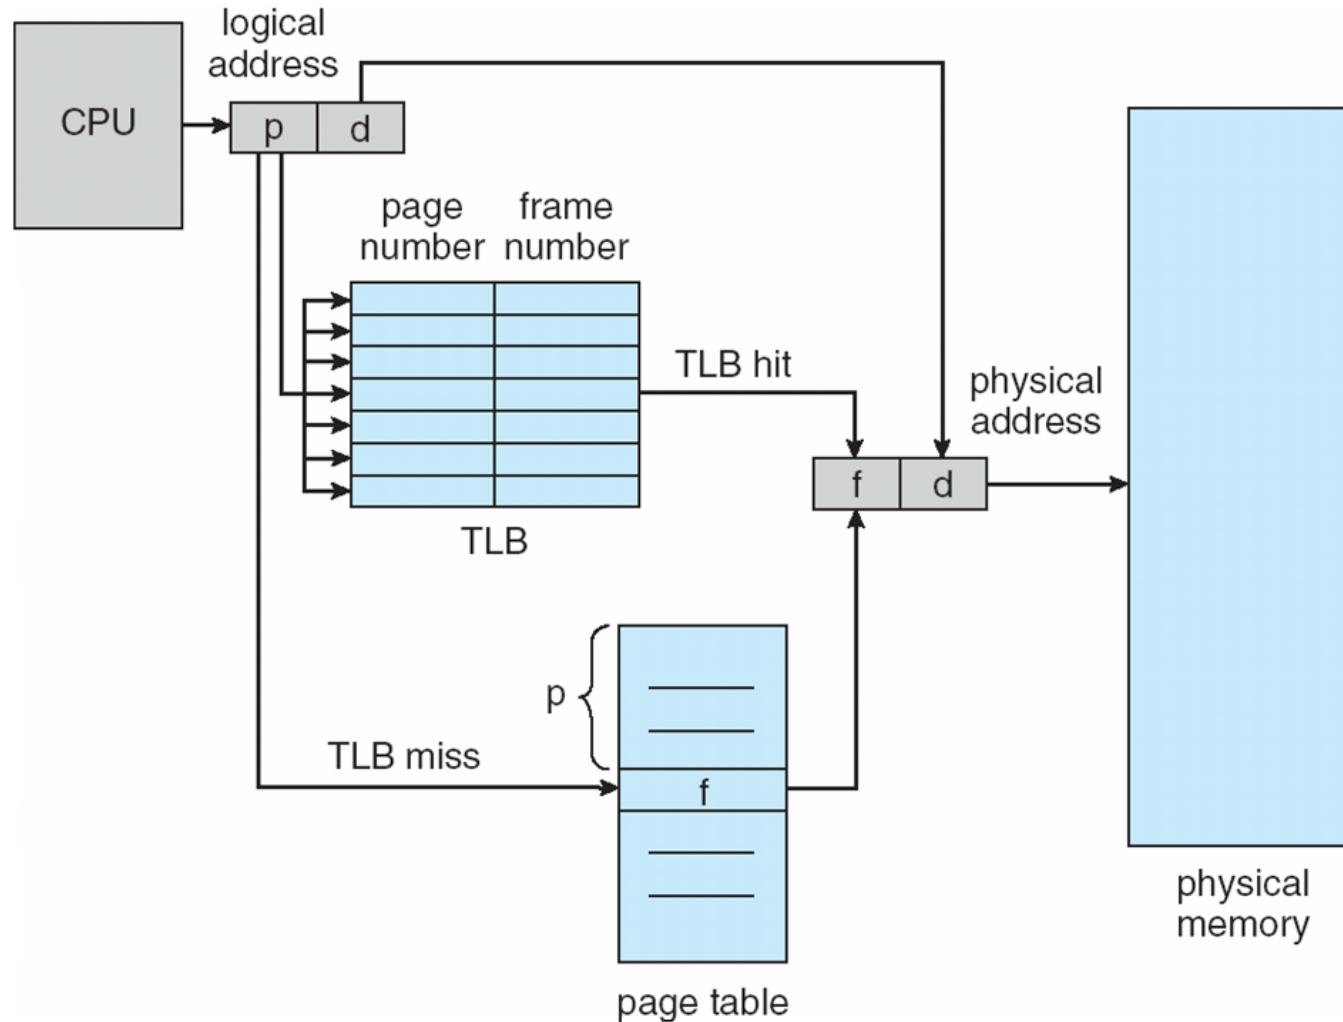
Page #	Frame #

- Address translation (p, d)
 - If p is in associative register, get Frame# out
 - Otherwise get Frame# from page table in memory





Paging Hardware With TLB





Effective Access Time

- Associative Lookup = ε time unit
 - Can be < 10% of memory access time
- Hit ratio = α (% of times that a page number is found in TLB)
- **Effective Access Time (EAT)**
$$\begin{aligned} EAT &= (1 + \varepsilon) \alpha + (2 + \varepsilon)(1 - \alpha) \\ &= 2 + \varepsilon - \alpha \end{aligned}$$

Even larger for multi-level page tables
- Consider $\alpha = 80\%$, $\varepsilon = 20\text{ns}$ for TLB search, 100ns for memory access
 - $EAT = 0.80 \times 100 + 0.20 \times 200 = 120\text{ns}$
- Consider more realistic hit ratio $\alpha = 99\%$, $\varepsilon = 20\text{ns}$ for TLB search, 100ns for memory access
 - $EAT = 0.99 \times 100 + 0.01 \times 200 = 101\text{ns}$





Structure of the Page Table

- Memory structures for paging can get huge using straight-forward methods
 - Consider a 32-bit logical address space as on modern computers: 2^{32} byte memory locations
 - Page size of 4 KB (4×2^{10} byte = 2^{12} byte)
 - Page table would have ~1 million entries ($2^{32} / 2^{12}$)
 - If each page table entry is 4 bytes
 - ▶ 4 MB ($2^{20} \times 4$ byte) of physical address space / memory for page table alone – equivalent to 1024 pages of 4KB each
 - ▶ **Don't want to allocate that contiguously in main memory**





Structure of the Page Table

- Memory structures for paging can get huge using straight-forward methods
 - Consider a 32-bit logical address space as on modern computers: **2^{32} byte memory locations**
 - Page size of 4 KB (4×2^{10} byte = 2^{12} byte)
 - **Page table** would have \sim **1 million entries** ($2^{32} / 2^{12}$)
 - If each page table entry is 4 bytes
 - ▶ 4 MB ($2^{20} \times 4$ byte) of physical address space / memory for page table alone – equivalent to 1024 pages of 4KB each
 - ▶ **Don't want to allocate that contiguously in main memory**

We moved to paging because we wanted to avoid allocating large blocks of contiguous memory

But we are back to the same problem – as logical address space increases, page tables need large memory blocks





Structure of the Page Table

- Memory structures for paging can get huge using straight-forward methods
 - Consider a 32-bit logical address space as on modern computers: 2^{32} byte memory locations
 - Page size of 4 KB (4×2^{10} byte = 2^{12} byte)
 - Page table would have ~1 million entries ($2^{32} / 2^{12}$)
 - If each page table entry is 4 bytes
 - ▶ 4 MB ($2^{20} \times 4$ byte) of physical address space / memory for page table alone – equivalent to 1024 pages of 4KB each
 - ▶ **Don't want to allocate that contiguously in main memory**

■ Solutions

- Hierarchical Paging
- Hashed Page Tables
- Inverted Page Tables





Hierarchical Page Tables

- Break up the logical address space into multiple page tables
- A simple technique is a two-level page table
- We then **page the page table**





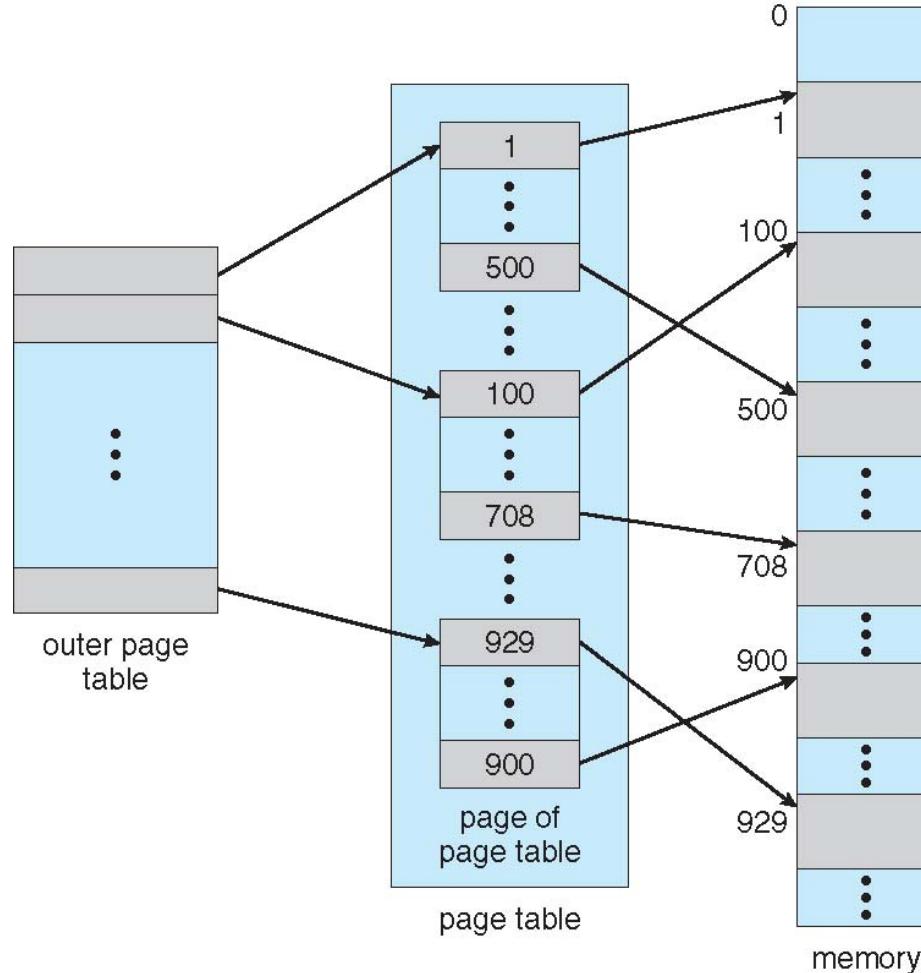
Hierarchical Page Tables

- Break up the logical address space into multiple page tables
- A simple technique is a two-level page table
- We then **page the page table (POPT: pages of page table)**





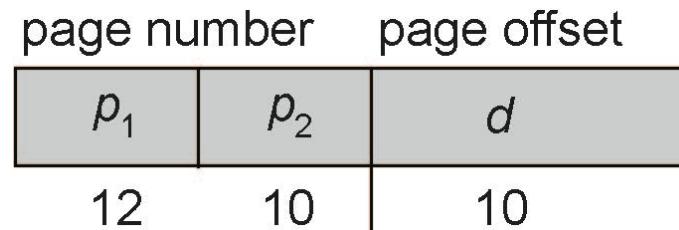
Two-Level Page-Table Scheme





Two-Level Paging Example

- A logical address (on 32-bit machine with 1K page size) is divided into:
 - a page number consisting of 22 bits
 - a page offset consisting of 10 bits
- Since the page table is paged, the page number is further divided into:
 - a 12-bit page number
 - a 10-bit page offset
- Thus, a logical address is as follows:

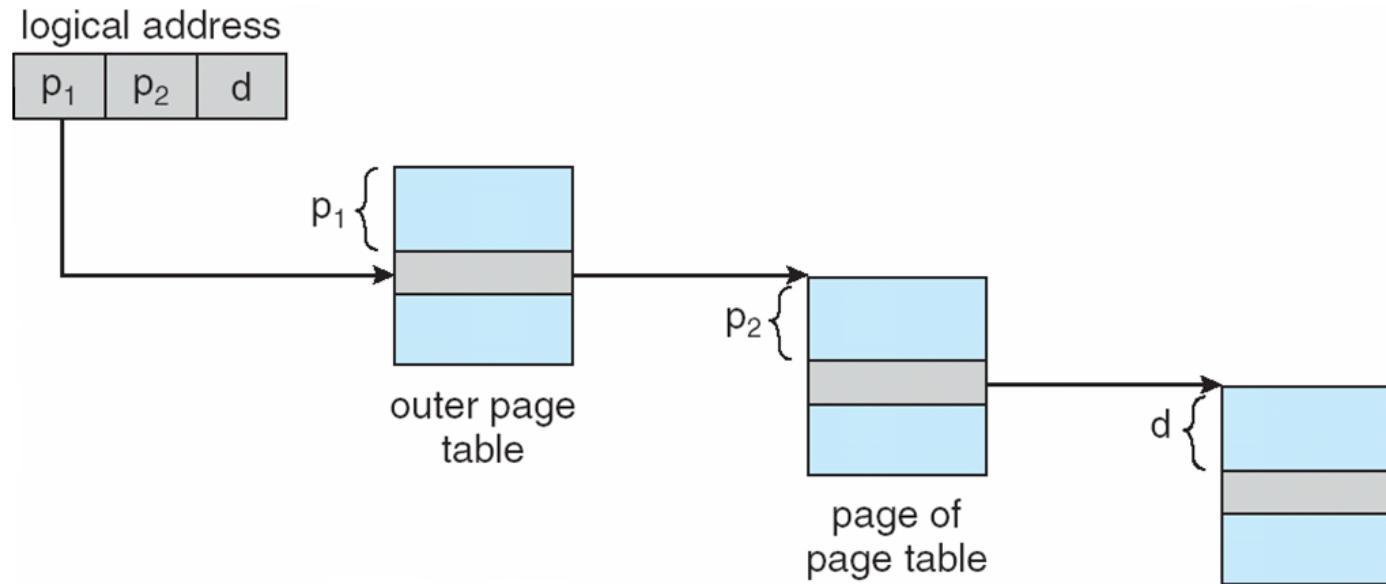


- where p_1 is an index into the outer page table, and p_2 is the displacement within the page of the inner page table



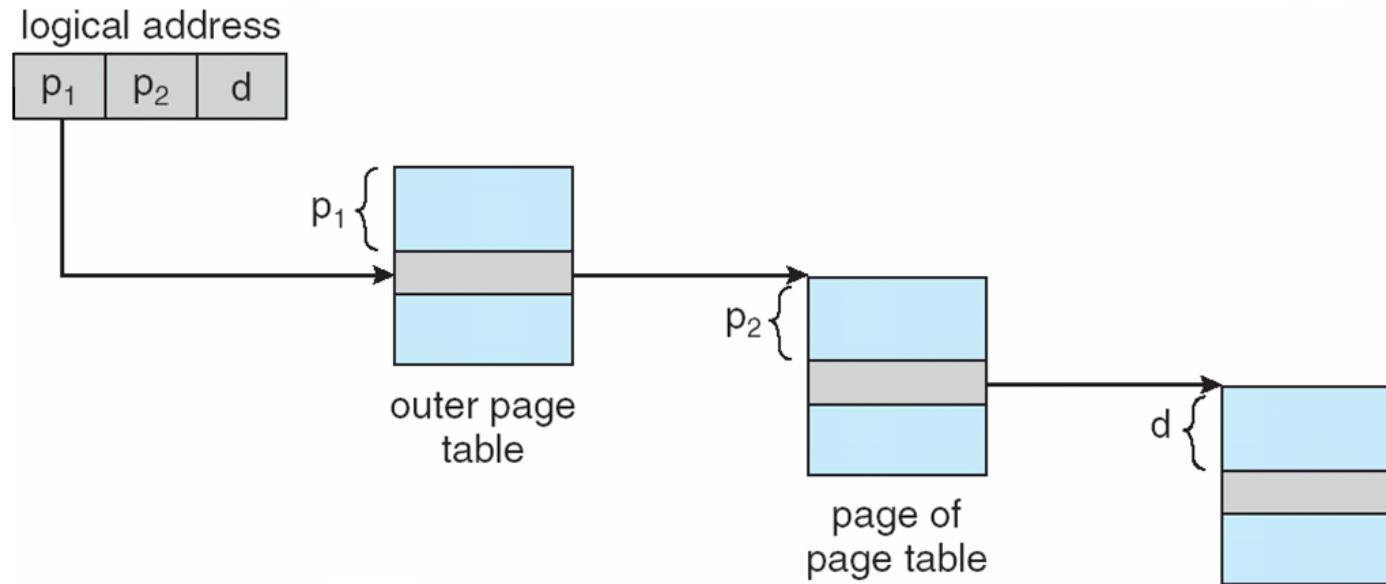


Address-Translation Scheme





Address-Translation Scheme



Cost: Three memory
Accesses
(previously two)



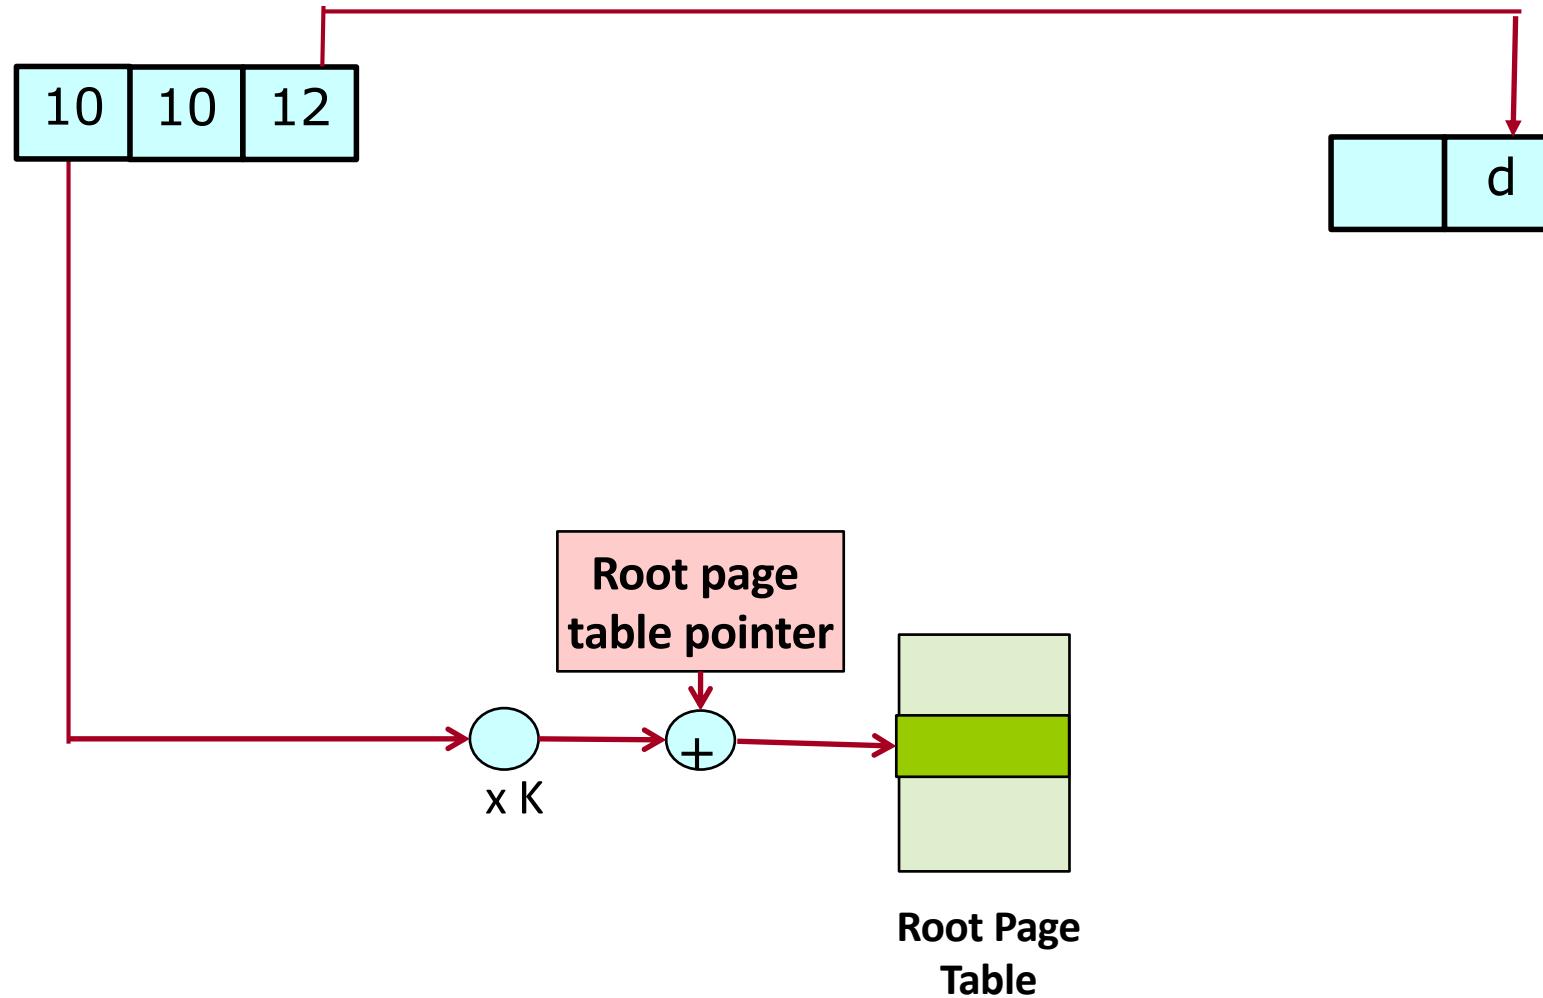


Implementation of 2-level paging



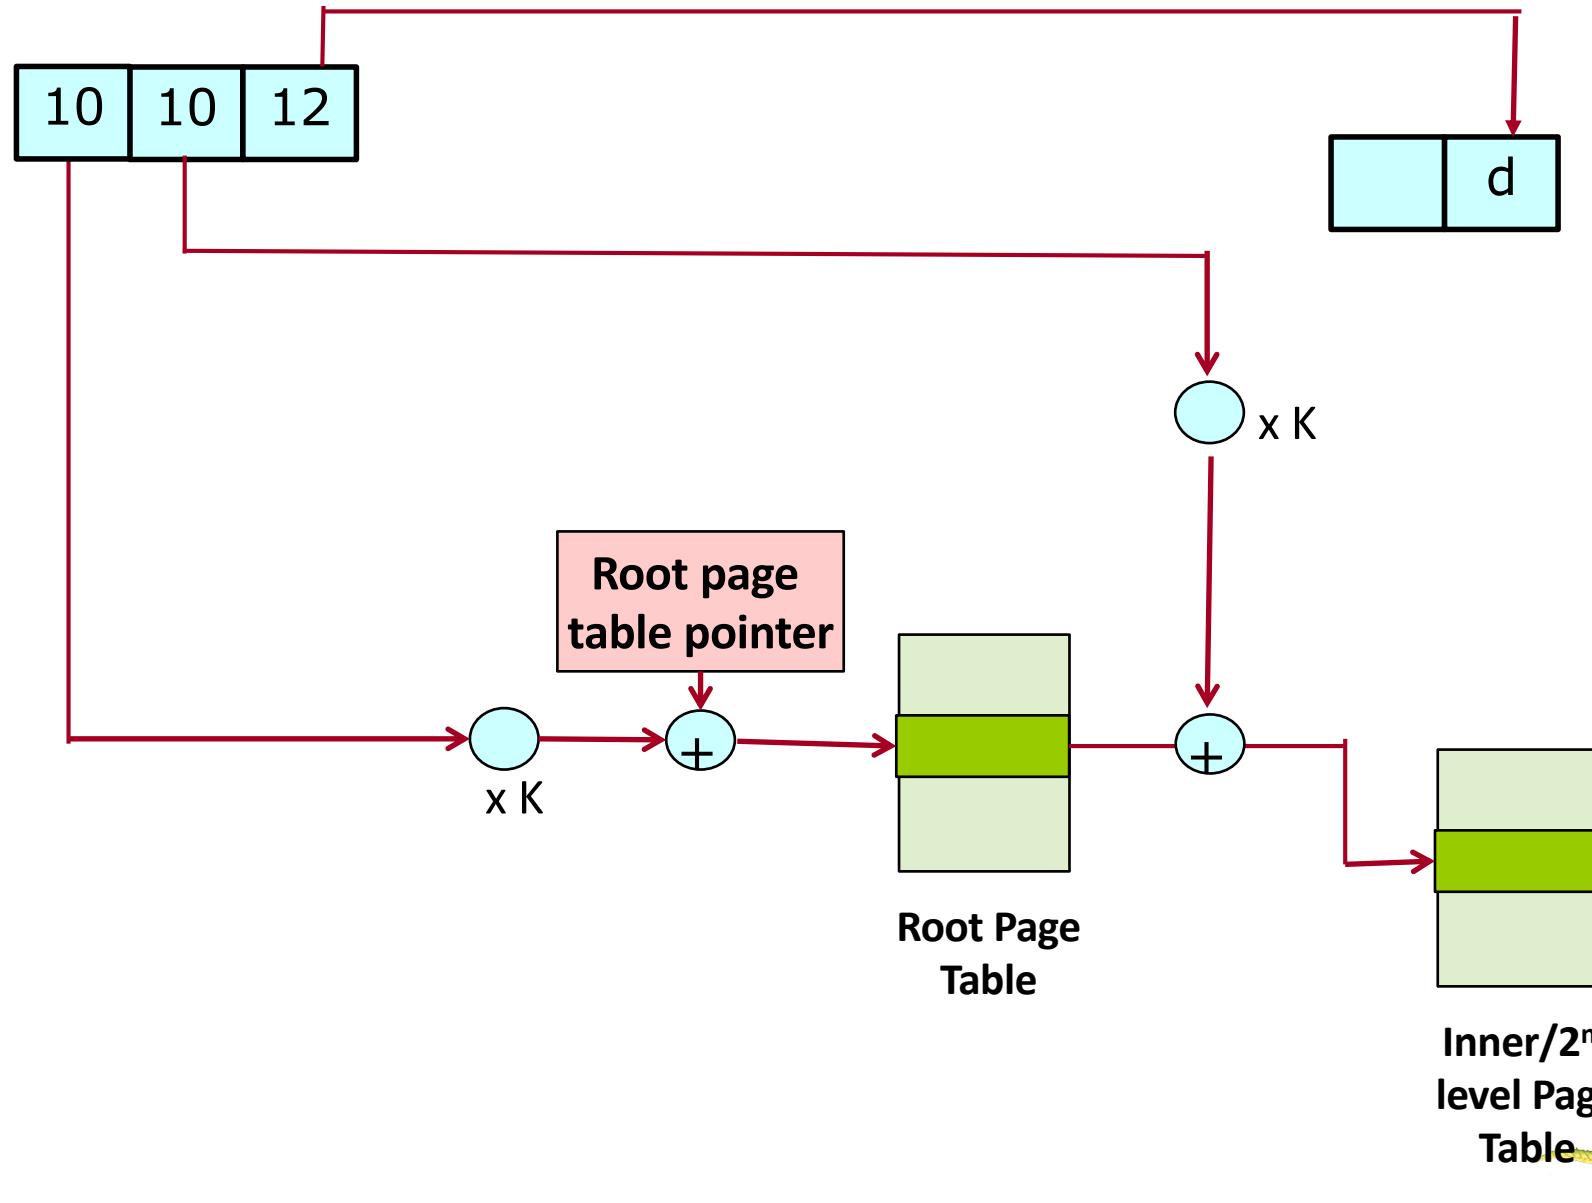


Implementation of 2-level paging



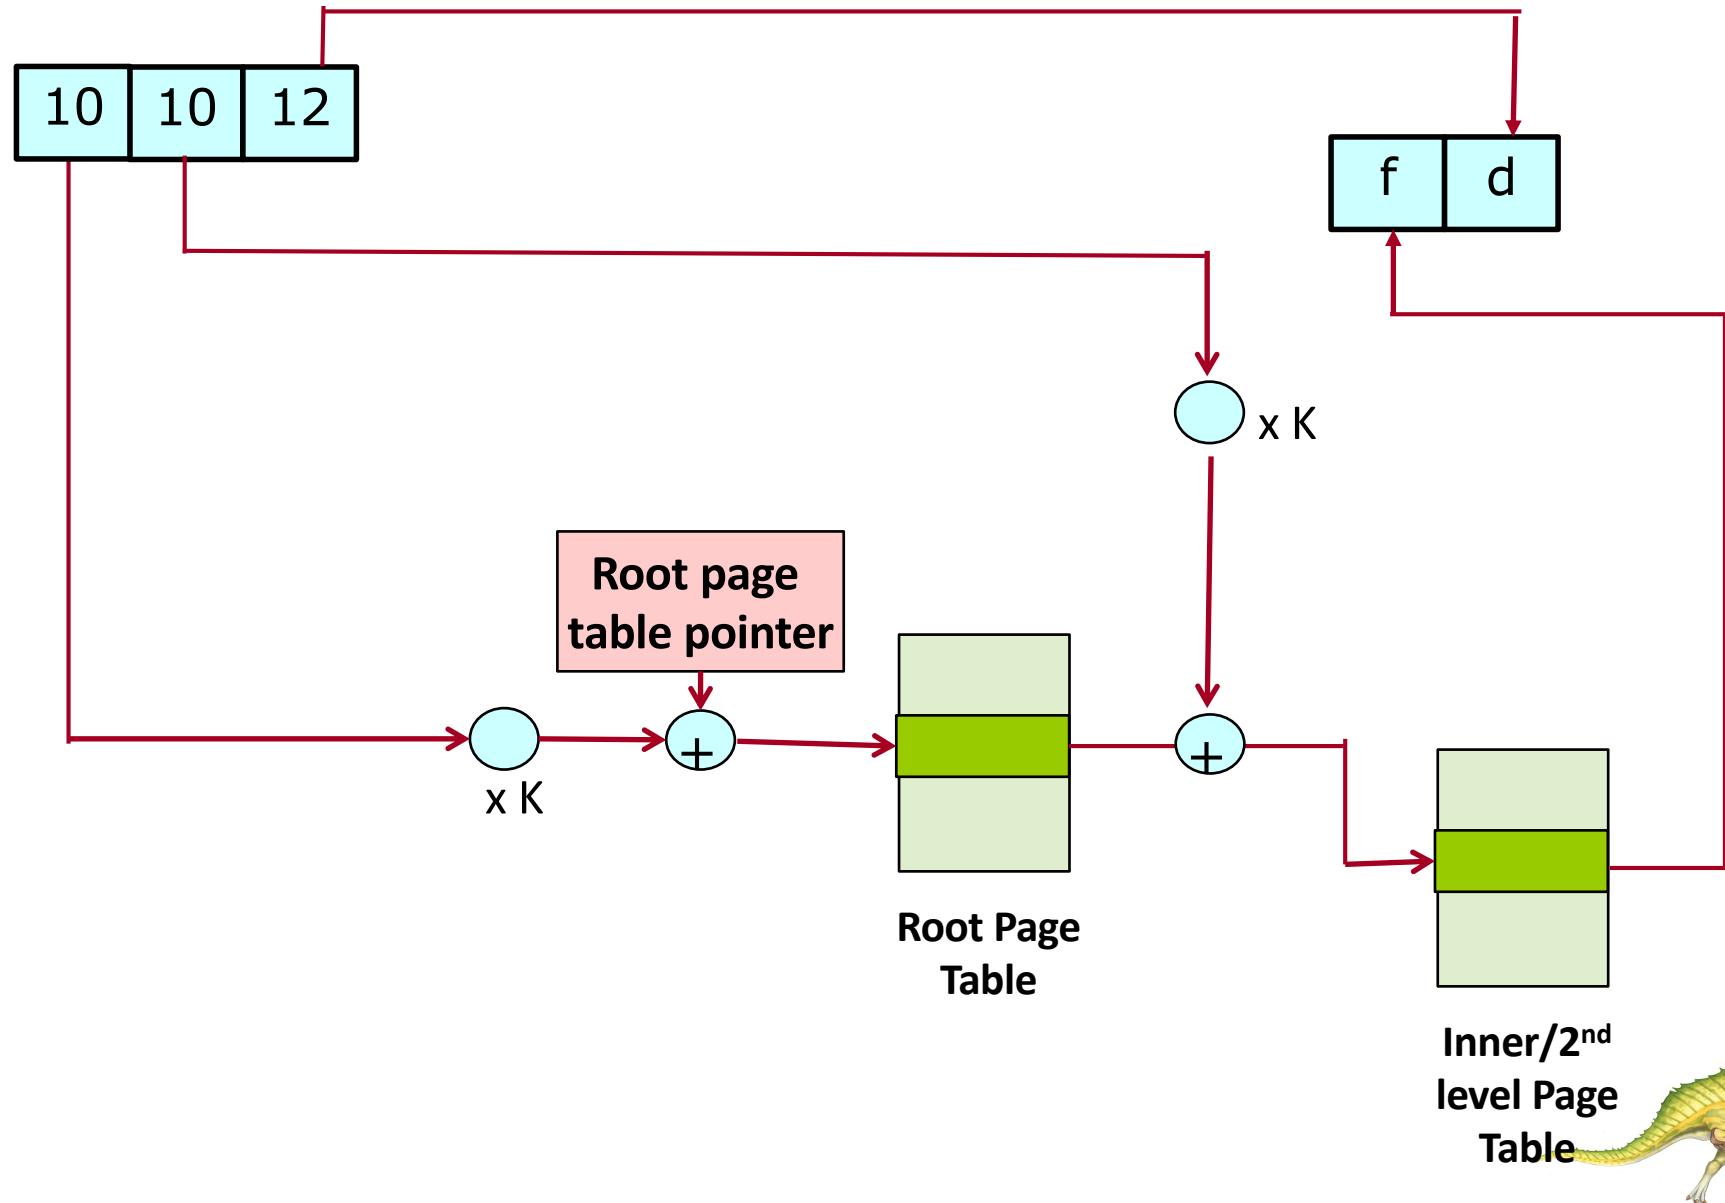


Implementation of 2-level paging





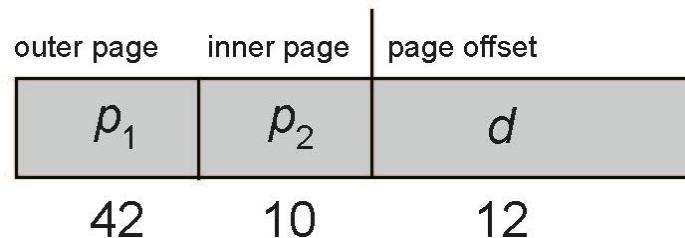
Implementation of 2-level paging





64-bit Logical Address Space

- Even two-level paging scheme not sufficient
- If page size is 4 KB (2^{12})
 - Then page table has 2^{52} entries
 - If 2-level scheme, inner/2nd level page tables could be 2^{10} 4-byte entries
 - Address would look like



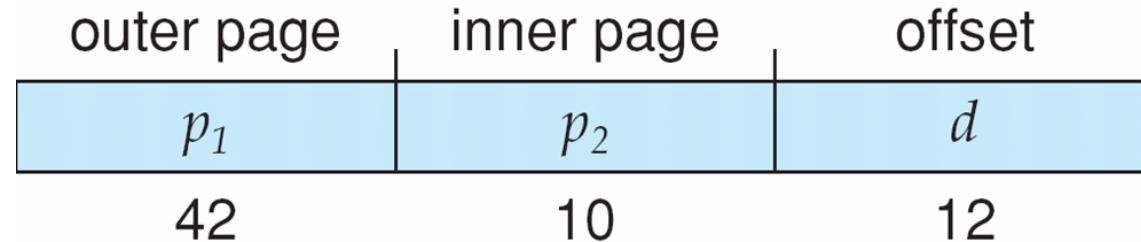
- Outer page table has 2^{42} entries or 2^{44} bytes = **2^{14} GB**

Won't work!





Three-level Paging Scheme



- One solution is to add a 2nd outer page table





Three-level Paging Scheme

outer page	inner page	offset
p_1	p_2	d
42	10	12

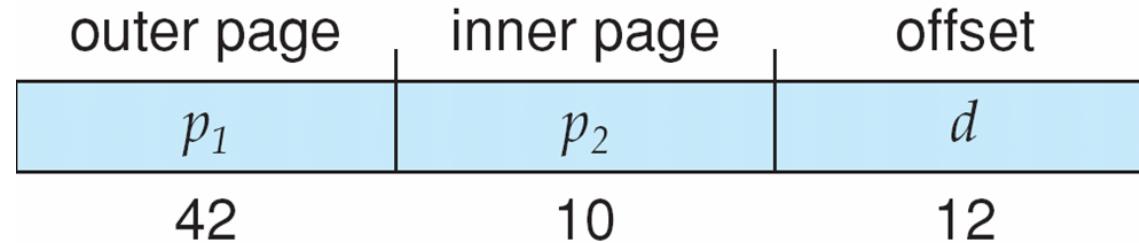
- One solution is to add a 2nd outer page table

2nd outer page	outer page	inner page	offset
p_1	p_2	p_3	d
32	10	10	12

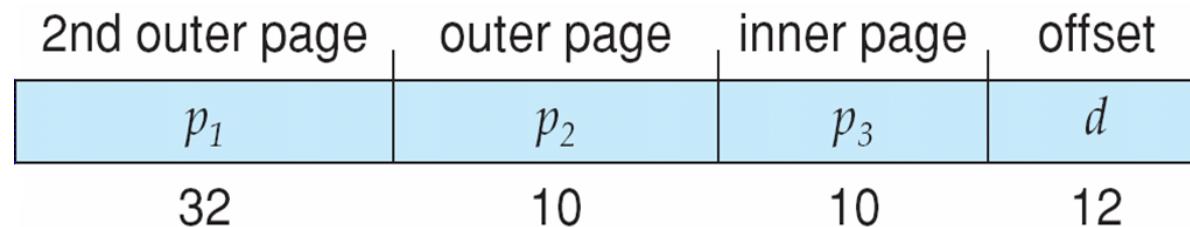




Three-level Paging Scheme



- One solution is to add a 2nd outer page table

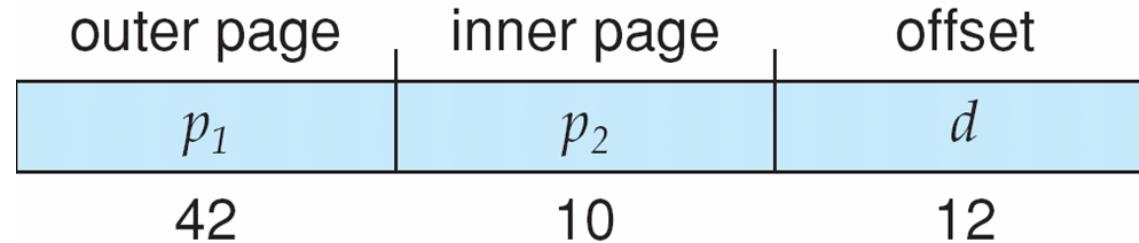


- But in this example the 2nd outer page table is still 2^{34} bytes in size
 - Even 16 GB is too huge to store contiguously
 - Also 4 memory access to get to one physical memory location

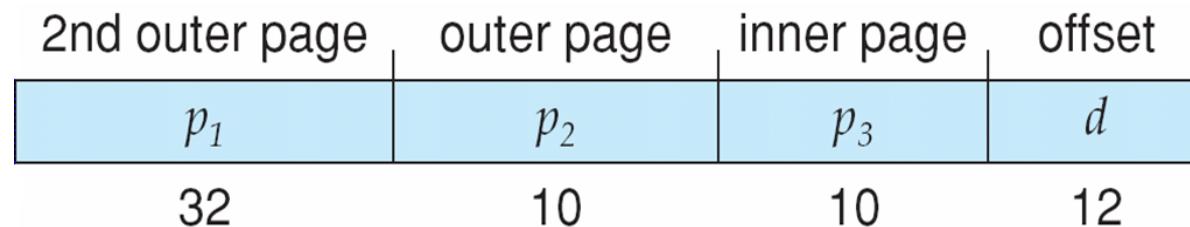




Three-level Paging Scheme



- One solution is to add a 2nd outer page table



- But in this example the 2nd outer page table is still 2^{34} bytes in size
 - Even 16 GB is too huge to store contiguously
 - Also 4 memory access to get to one physical memory location

Multi-level paging schemes do not scale well
Need alternate mechanisms





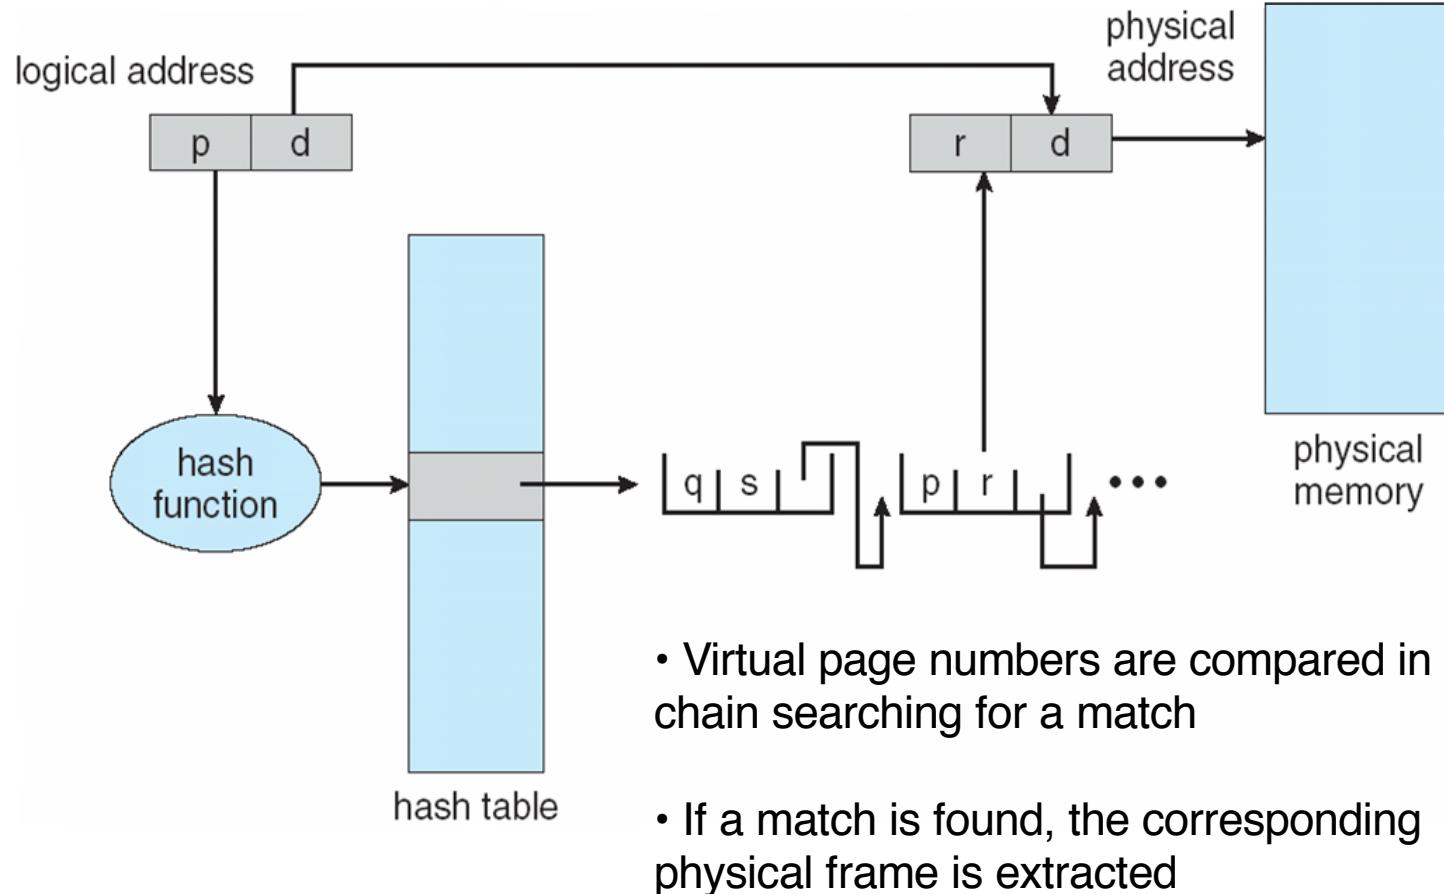
Hashed Page Tables

- Common in address spaces > 32 bits
- The virtual page number is hashed into a page table
 - This page table contains a chain of elements hashing to the same location
- Each element contains
 - (1) the virtual page number
 - (2) the value of the mapped physical page frame
 - (3) a pointer to the next element





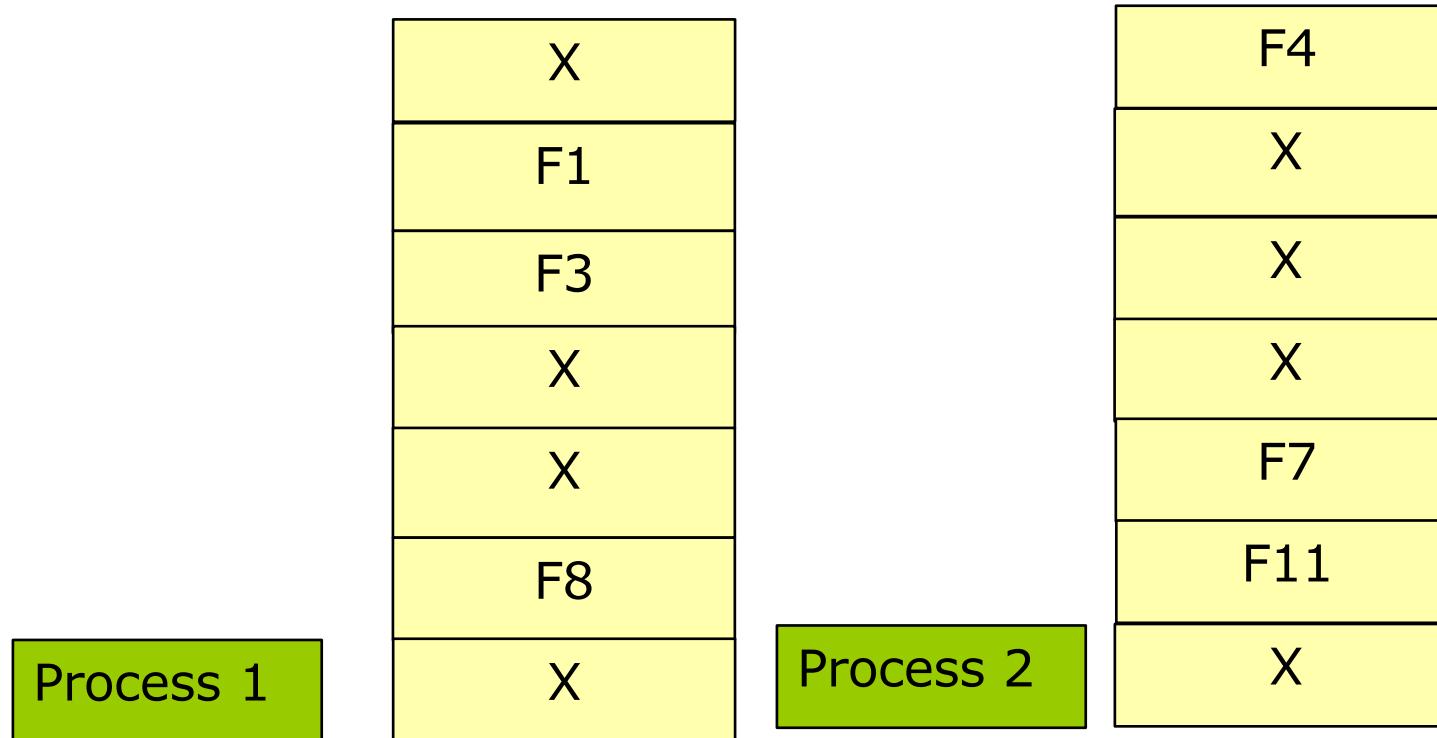
Hashed Page Table





Inverted Page Table: Background

- Till now what we have discussed:
 - Each process has a page table that keeps track of all logical pages of this process





Inverted Page Table

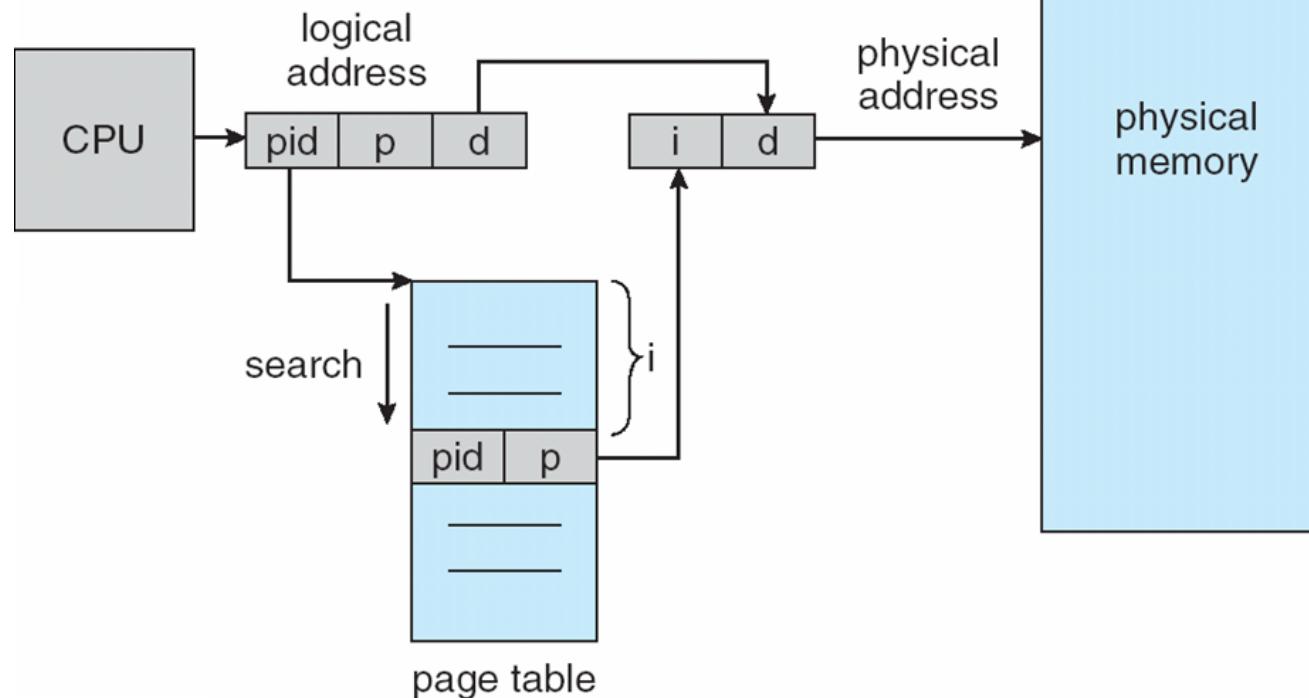
- Rather than each process having a page table and keeping track of all possible logical pages, track all physical pages (frames)
- Inverted page table: One entry for each real page (frame) of memory
 - Entry = the virtual address of the page stored in that real memory location + information of the process that owns the page





Inverted Page Table Architecture

A logical address is a triple:
<process-id, page-number, offset>





Inverted Page Table

- Decreases memory needed to store each page table, but **increases time needed to search the table for a page reference**
 - Searching uses the CPU-generated virtual address, but the inverted page table is sorted by the physical address

- Can make searching physical address relatively fast:
 - Use hash table (as described earlier)
 - TLB can accelerate access further (even before consulting hash table)





Summary

- Contiguous allocation with Relocation register
- Paging – avoid external fragmentation, process sees contiguous memory space, but physically non-contiguous
- Where to store the Page Table? In the memory
 - Cost – multiple memory accesses for each data access
 - Optimization for faster access – use TLB (associative cache)
- Structure of the Page Table in memory
 - Page Table can be too large for contiguous memory allocation
 - Hierarchical multi-level Page Table
 - ▶ Store pages of page table (POPT)
 - ▶ Does not scale well as logical address space becomes larger
 - Hashed Page Table
 - Inverted Page Table





What a programmer wants from memory

- Memory addresses will start from 0
- Memory will be contiguous
- I should be able to dynamically increase the memory
- I should have as much memory as I want

First 3 requirements satisfied with paging. Also, we now have ways of allowing huge memory to processes.

I may even want more memory than the physical mem size. I do not want to be limited to the physical memory size on a particular system.

Virtual memory – to be discussed next

